

AN IMAGE IS WORTH 1000 LIES: ADVERSARIAL TRANSFERABILITY ACROSS PROMPTS ON VISION-LANGUAGE MODELS

Haochen Luo*, Jindong Gu*†, Fengyuan Liu, Philip Torr

Torr Vision Group, University of Oxford

Parks Road, Oxford OX1 3PJ, UK

haochen.luo@outlook.com, jindong.gu@outlook.com

ABSTRACT

Different from traditional task-specific vision models, recent large VLMs can readily adapt to different vision tasks by simply using different textual instructions, i.e., prompts. However, a well-known concern about traditional task-specific vision models is that they can be misled by imperceptible adversarial perturbations. Furthermore, the concern is exacerbated by the phenomenon that the same adversarial perturbations can fool different task-specific models. Given that VLMs rely on prompts to adapt to different tasks, an intriguing question emerges: Can a single adversarial image mislead all predictions of VLMs when a thousand different prompts are given? This question essentially introduces a novel perspective on adversarial transferability: cross-prompt adversarial transferability. In this work, we propose the Cross-Prompt Attack (CroPA). This proposed method updates the visual adversarial perturbation with learnable prompts, which are designed to counteract the misleading effects of the adversarial image. By doing this, CroPA significantly improves the transferability of adversarial examples across prompts. Extensive experiments are conducted to verify the strong cross-prompt adversarial transferability of CroPA with prevalent VLMs including Flamingo, BLIP-2, and InstructBLIP in various different tasks. Our source code is available at <https://github.com/Haochen-Luo/CroPA>.

1 INTRODUCTION

Previously task-specific vision models have demonstrated remarkable capabilities in various visual tasks such as image classification (He et al., 2016) and image captioning (Yao et al., 2018; Yang et al., 2019). These models are designed to extract specific information which is pre-defined during the construction phase of the model. Recently, large Vision-Language Models (VLMs) (Gu et al., 2023a; Li et al., 2022; Alayrac et al., 2022; Li et al., 2023; Zhu et al., 2023) have emerged, providing a more unified approach to addressing computer vision tasks. Instead of relying merely on the image input, VLMs integrate information from both images and associated textual prompts, enabling them to perform varied vision-related tasks by utilizing appropriate prompts. This versatility paves the way for exploring various visual tasks.

Those task-specific models are known to be vulnerable to adversarial examples (Goodfellow et al., 2014; Szegedy et al., 2013; Wu et al., 2022; Gu et al., 2022). These examples are developed by adding subtle perturbations, typically invisible to humans, to original samples. These perturbations can significantly degrade the performance of such models significantly (Madry et al., 2017). Compounding the mentioned vulnerabilities, the adversarial examples further exhibit transferability across models (Gu et al., 2023b). This implies that these adversarial examples can attack target models beyond those for which they were specifically crafted (Liu et al., 2016; Tramèr et al., 2017).

Given that adversarial examples have transferability across tasks (Salzmann et al., 2021; Gu et al., 2023b), an interesting question emerges: **Is an adversarial example transferable across prompts?** We

*Equal contribution, †Corresponding author

dub it cross-prompt adversarial transferability, which means that regardless of the prompts provided, the model output can consistently be misled by an adversarial example. For example, if the target text is set to “unknown”, a model, deceived by an adversarial example exhibiting cross-prompt transferability, will always predict the word “unknown” regardless of the prompts. In this case, VLMs are incapable of extracting information from the image, even when different textual prompts are presented.

Exploring cross-prompt adversarial examples is crucial to revealing the prompt-related vulnerability of VLM and protecting image information. From the perspective of an attacker, adversarial examples with high cross-prompt transferability can mislead large VLMs to generate malicious outputs even when queried with various benign prompt questions. From the defensive perspective, the potential to obfuscate image information through human-imperceptible perturbations can cause the model to uniformly output a predefined target text. This can prevent malicious usage of large VLMs for unauthorized extraction of sensitive information from personal images.

Our experiments have shown the cross-prompt transferability created with a single prompt is highly limited. An intuitive approach to increase the cross-prompt transferability is to use multiple prompts during its creation stage. However, the improvement in cross-prompt transferability of these baseline approaches converges quickly with the increase in prompts. To further improve the cross-prompt transferability, we proposed Cross-Prompt Attack (CroPA), which creates more transferable adversarial images by utilising the learnable prompts. These prompts are optimised in the opposite direction of the adversarial image to cover more prompt embedding space.

In the experiments, prompts for three popular vision-language tasks are used, including image classification, image captioning, and visual question answering (VQA). We examined the effectiveness of our approach on three prevalent VLMs, Flamingo (Alayrac et al., 2022), BLIP-2 (Li et al., 2023) and InstructBLIP (Dai et al., 2023). Experimental results have demonstrated that CroPA consistently outperforms the baseline methods under different settings with different attack targets.

Our contributions can be summarized as follows:

- We introduce cross-prompt adversarial transferability, an important perspective of adversarial transferability, contributing to the existing body of knowledge on VLMs’ vulnerabilities.
- We propose a novel algorithm Cross-Prompt Attack (CroPA), designed to enhance cross-prompt adversarial transferability.
- Extensive experiments are conducted to verify the effectiveness of our approach on various VLMs and tasks. Moreover, we provide further analysis to understand our approach.

2 RELATED WORK

Adversarial transferability Foundational studies by (Szegedy et al., 2013; Goodfellow et al., 2014) unveil the property of neural networks to misclassify images by adding seemingly imperceptible adversarial perturbations to the inputs. The created adversarial samples can also fool unseen models (Gu et al., 2023b; Yu et al., 2023; Liu et al., 2016; Papernot et al., 2016). Besides, Mopuri et al. (2017); Moosavi-Dezfooli et al. (2017) shows that an adversarial perturbation can be still deceptive when added to different images. Beyond models and images, the domain of adversarial transferability extends its reach to different tasks (Naseer et al., 2018; 2019; Lu et al., 2020; Salzmann et al., 2021). For example, adversarial examples designed to attack image classification systems are not limited in their scope but also fail other tasks, such as object detection. In light of the revealed transferability of adversarial examples across models, images, and tasks, the recent advancements in vision-language models introduce a new dimension to be explored. Specifically, this work delves into the transferability across textual prompts within the realm of VLMs.

Adversarial Robustness of Vision-Language Models The majority of prior research on adversarial attacks on vision-language models are mostly task-specific attacks. For example, there is a series of works to manipulate the model output in image captioning tasks (Xu et al., 2019; Zhang et al., 2020; Aafaq et al., 2021; Chen et al., 2017a). Similarly, in visual question answering, works such as Fooling VQA (Xu et al., 2018; Kaushik et al., 2021; Kovatchev et al., 2022; Li et al., 2021; Sheng et al., 2021; Zhang et al., 2022a) mislead the attention region in object detectors to affect the model output. Nevertheless, the vision-language models used in these methods are highly task-specific, utilizing lightweight CNN-RNN architectures that lack the capability for in-context learning. Con-

sequently, adapting these methods to contemporary VLMs gives challenges. There are recent works on the adversarial robustness of large VLMs that consider the adversarial attack from the vision modality. Concretely, Zhao et al. (2023) explored the adversarial robustness of recent large vision-language models such as BLIP (Li et al., 2022) and BLIP-2 (Li et al., 2023) under the black box setting including query-based and transfer-based methods to craft adversarial examples. Instead of transferability across models, this work introduces cross-transferability.

3 APPROACH

In this section, we first describe the concept of cross-prompt adversarial transferability. We then present the baseline approach utilising one or multiple prompts to craft adversarial examples, and the CroPA method which incorporates learnable prompt to enhance cross-prompt transferability.

3.1 PROBLEM FORMULATION

Consider x_v to be a clean image without perturbations induced and let x_t denote a prompt. The function f represents a VLM. The term δ_v signifies the visual perturbation added to the image x_v and is bound by the constraints $\|\delta_v\|_p \leq \epsilon_v$, where ϵ_v is the image perturbation magnitude.

- **Targeted Attack:** In a targeted attack, the objective is to generate a visual perturbation, denoted as δ_v , which when applied to the original input x_v , creates an adversarial example $x_v + \delta_v$. **This adversarial example is structured to mislead the model into producing a predefined targeted output text T , regardless of the given prompt.**
- **Non-Targeted Attack:** Contrarily, in a non-targeted attack, the adversarial example is crafted not to lead the model to a specific predefined output but rather to any incorrect output. **The goal here is to ensure that the model’s output, when fed with the adversarial example, diverges from the output generated with a clean, unaltered image as input.**

Attack success rate (ASR) is the evaluation metric for cross-prompt transferability, which is defined as the ratio of the number of successful attacks to the total number of attacks. For the targeted attack, the attack is considered to be successful only if the prediction exactly matches our target text. For non-targeted attacks, the attack is successful if the model is misled to generate the text different from the prediction with the clean image.

3.2 BASELINE APPROACH

To generate adversarial examples for VLMs, an image perturbation can be optimized based on a single prompt; this method is referred to as **Single-P**. To enhance the cross-prompt transferability of the perturbations, a straightforward approach is to utilize multiple prompts while updating the image perturbation, a method denoted by **Multi-P**. The algorithms of Single-P and Multi-P are detailed below.

Let $\mathcal{X}_t = \{x_t^1, x_t^2, \dots, x_t^k\}$ represent a collection of textual prompt instances. The ultimate goal is to derive a visual perturbation, δ_v , ensuring that for every instance from \mathcal{X}_t , the model yields either the predefined target text T in the targeted attack setting, or text deviating from the original output in a non-targeted attack setting.

The optimization objectives for targeted and non-targeted settings are formulated as follows: **For the targeted attack, the objective is to minimize the language modelling loss, \mathcal{L} , associated with generating the target text T . This optimization can be mathematically represented as:**

$$\min_{\delta_v} \sum_{i=1}^k \mathcal{L}(f(x_v + \delta_v, x_t^i), T) \quad (1)$$

Here, the goal is to alter the input subtly to mislead the model into producing the predefined text T across various prompt instances, effectively minimizing the discrepancy between the model’s output and the target text.

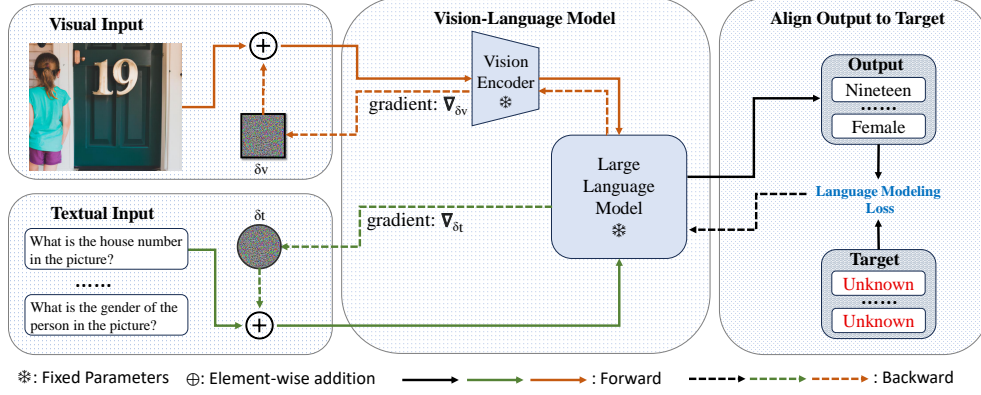


Figure 1: Overview of CroPA’s framework under the targeted attack setting. Both the image perturbation δ_v and the prompt perturbation δ_t are learnable but the prompt perturbation does not collaborate with δ_v to deceive the model. They are optimised with the opposite goals: δ_v aims to minimise the language modelling loss while the δ_t aims to maximise the language modelling loss. The update frequency of the image perturbation and prompt perturbation can be different.

For the non-targeted attack, the objective is to maximize the language modelling loss \mathcal{L} , between the text produced by the model with adversarial examples and clean images: $\max_{\delta_v} \sum_{i=1}^k \mathcal{L}(f(x_v + \delta_v, x_t^i), f(x_v, x_t^i))$. The aim here is not to guide the model to a specific output but to any output diverging from what would have been produced with an unaltered input, emphasizing the maximization of the discrepancy in the model’s responses.

3.3 CROSS-PROMPT ATTACK (CROPA)

In order to create an adversarial image with stronger cross-prompt transferability, we propose an algorithm termed Cross-Prompt Attack (**CroPA**). The baseline approach constrains prompts to decoded text representations, namely the fixed hard prompt. In the CroPA framework, we not only used varied numbers of prompts but also introduced learnable prompt perturbation for the prompt embedding during the optimisation phase.

In Figure 1, we provide the illustration of the CroPA framework through an example, where the image perturbation is optimised with the target text “unknown” to hide sensitive information such as address and gender. The image perturbation δ_v is optimised to minimise the loss of generating the target text “unknown”, while the prompt perturbation δ_t is updated in the opposite direction to maximise the loss of generating the target text. The prompt perturbation gains increasingly stronger cross-prompt transferability during this competitive update process.

Concretely, the optimisation steps of the adversarial image are detailed in Algorithm 1. In the beginning, the image adversarial perturbation and the prompt adversarial perturbation are randomly initialised. During the forward pass, these two perturbations are added to the clean image and the clean prompt embedding respectively. After the forward pass, the gradient of language modelling loss with respect to the image and text can be obtained through backward propagation. While the image perturbation is updated with gradient descent to minimize the language modelling loss between the prediction and the target sentence, the adversarial prompt is updated with gradient ascent to maximize the loss. The parameters α_1 and α_2 are the updated step sizes for the adversarial image and adversarial prompts. The optimisation algorithm is PGD (Madry et al., 2017) with the L-infinity norm being the specific norm used in the experiments.

The update of the adversarial image and the adversarial text can be viewed as a min-max process. Considering a vision-language model f that takes an image x_v and a text prompt x_t as input, the objective is to obtain the perturbations δ_v for x_v that minimises the language modeling loss \mathcal{L} of generating the targeted sentence, and the perturbations δ_t for x_t that maximises the loss. For the targeted attack, the optimisation of the formula can be written as:

$$\min_{\delta_v} \max_{\delta_t} \mathcal{L}(f(x_v + \delta_v, x_t + \delta_t), T) \quad (2)$$

Algorithm 1 CroPA: Cross Prompt Attack

Require: Model f , Target Text T , vision input x_v , prompt set X_t , perturbation size ϵ , step size of perturbation updating α_1 and α_2 , number of iteration steps K , adversarial prompt update interval N

Ensure: Adversarial example x'_v

```

1: Initialise  $x'_v = x_v$ 
2: for step = 1 to  $K$  do
3:   Uniformly sample the prompt  $x_t^i$  from  $\mathcal{X}_t$ 
4:   if  $x_t^{i'}$  is not initialised then
5:     Initialise  $x_t^{i'} = x_t^i$ 
6:   end if
7:   Compute gradient for adversarial image :  $g_v = \nabla_{x_v} \mathcal{L}(f(x'_v, x_t^i), T)$ 
8:   Update with gradient descent:  $x'_v = x'_v - \alpha_1 \cdot \text{sign}(g_v)$ 
9:   if mod(step,  $N$ ) == 0 then
10:    Compute gradient for adversarial prompt:  $g_t = \nabla_{x_t} \mathcal{L}(f(x'_v, x_t^i), T)$ 
11:    Update with gradient ascent:  $x_t^{i'} = x_t^{i'} + \alpha_2 \cdot \text{sign}(g_t)$ 
12:   end if
13:   Project  $x'_v$  to be within the  $\epsilon$ -ball of  $x_v$ :  $x'_v = \text{Clip}_{x_v, \epsilon}(x'_v)$ 
14: end for
15: return  $x'_v$ 

```

Similarly, for the non-targeted attack, the optimisation can be expressed as: $\max_{\delta_v} \min_{\delta_t} \mathcal{L}(f(x_v + \delta_v, x_t + \delta_t), f(x_v, x_t))$. The visual perturbation δ_v is optimised to maximise the loss of generating $f(x_v, x_t)$ so that the model is deceived to generate the output different from the original one. Prompt perturbation δ_t is optimised to minimise the language modelling loss of generating $f(x_v, x_t)$. For both targeted attack and non-targeted attack, the image perturbation is clipped to the ϵ to ensure the invisibility of the image perturbation. We use the parameter N denoting the update interval to control the update frequency of image perturbation and prompt perturbation: the image perturbation for N times, and the prompt perturbation updates once. Please note that the prompt perturbations are added only during the optimisation phase and they are not added during the testing phase.

4 EXPERIMENTS

Experimental Settings The dataset consists of both images and prompts. The images are collected from the validation dataset of MS-COCO datasets (Lin et al., 2014). The prompts for VQA consist of questions both agnostic and specific to the image content, which are referred as to VQA_{general} and VQA_{specific} in the following sections. The image-specific questions derive from the VQA-v2 (Goyal et al., 2017). We craft prompts for the questions agnostic to image content, image classification, and image captioning with diverse lengths and semantics. By default, the experiments are targeted attacks with the target text set to “unknown” to avoid the inclusion of high-frequency responses in vision-language tasks. Adversarial examples are optimised and tested under 0-shot settings. The number of prompts for Multi-P and CroPA is set to ten. Detailed prompts can be found in Appendix B. The VLMs used are Flamingo, BLIP-2, and InstructBLIP. We adopt the open-source OpenFlamingo-9B (Awadalla et al., 2023) for Flamingo. Attack Success Rate is used as a metric in our experiments. All the ASR scores reported in the following sections are averaged over three runs. The perturbation size is set to 16/255.

4.1 CROSS-PROMPT TRANSFERABILITY COMPARISON

The cross-prompt adversarial transferability is expected to be stronger if more prompts are given during the optimisation stage. To verify this assumption, we sample different numbers of prompts: 1, 5, 10, 50, and 100, and test the targeted ASR. The overall performance of the ASR of the baseline methods and CroPA with different numbers of prompts tested with Flamingo, BLIP-2 and InstructBLIP are shown in Figure 2.

For all the experiments conducted in this section, we selected the target text “unknown” to avoid the inclusion of high-frequency responses commonly found in vision-language tasks. Based on the experimental results, we can conclude the following points: 1) CroPA consistently outperforms the baseline approach for all models. As Figure 2 shows, CroPA achieves the best overall performance in all testing models for all different prompt numbers. CroPA also achieves the best individual performance in most tasks. The detailed data can be found in the supplementary. 2) More prompts




Figure 2: The targeted ASR of three methods tested on (a) Flamingo, (b) BLIP-2, and (c) InstructBLIP. Different numbers of prompts 1, 5, 10, 50, and 100 are used in the transferability test. Our CroPA achieve better cross-prompt adversarial transferability than Single-P and Multi-P.

increase the transferability, but convergence occurs rapidly. We can observe that in general more prompts increase the targeted ASR, especially when the number of prompts increases from one to five. However, starting from ten, the increase of cross-prompt transferability brought by more prompts becomes marginal, especially for the baseline approach. This pattern indicates that by adding more prompts, the baseline approach cannot surpass the performance of CroPA methods.

We also conducted experiments with cross-prompt transferability tests combined with other dimensions of transferability, models or images. Our experimental results have shown that adversarial transferability is limited. When tested across different prompts and images, the ASRs are all near 0. The ASRs over the models with significantly different architecture are near 0. For the pair of InstructBLIP and BLIP-2, an overall ASR of around 10% is achieved and the detailed data can be found in Table 8 in Appendix C. It should be noted that the language models used in BLIP-2 and InstructBLIP are two different models: OPT-2.7b (Zhang et al., 2022b) and Vicuna-7b (Zheng et al., 2023) respectively. Though the performances when combined with other perspectives of the transferability are not as strong as cross-prompt transferability alone, CroPA still consistently outperforms the baseline methods. Moreover, existing methods for enhancing cross-model and cross-image transferability are orthogonal to the methods proposed in this study. For example, the cross-image transferability can be enhanced by computing the perturbation over a larger number of images (Moosavi-Dezfooli et al., 2017). We extend the work of creating adversarial images with combined perspectives of transferability to future work.

4.2 CROPA WITH DIFFERENT TARGET TEXTS

To demonstrate that the effectiveness of the proposed CroPA methods is not constrained to the specific case of the target text “unknown”, we extend our evaluation to various other target texts. The experiment includes a selection of text with varied length and usage frequency. Both common expressions such as “very good” and less common words such as “metaphor” are tested.

From Table 1, we can observe that CroPA performs the best overall and in each individual task under different target text, though the ASR differs for different target prompts. This table also demonstrated that cross-prompt transferability is not associated with the semantic meaning or the word frequency of the target text. Semantically “not sure” is a reasonable generation result of VLMs to indicate the uncertainty of the response. However, its targeted ASR is much lower than “metaphor”, which is less commonly used as a response to the questions or instructions about an image. The results also show that the cross-prompt transferability is not determined by the target text length. The ASR of the three-word target text is not necessarily to be lower than the ASR of a one-word target text.

We also conducted the experiments by setting the target texts to harmful instructions such as “suicide”, “bomb” and “kidnap”. As shown in Table 9, the high ASR values indicate that the cross-prompt adversarial transferability can lead to security issues in real-world applications. From this perspective, the cross-prompt transferability is also related to the jailbreak Zou et al. (2023) for vision-language models, which also aims to deceive the model to generate harmful instructions.

In summary, the CroPA has been proven to be a stronger method than the baseline approach with different target texts. While different target texts do affect the cross-prompt adversarial transferability, it is unlikely that this is associated with the semantics or length of the prompts themselves.

Table 1: Targeted ASRs tested on Flamingo with different target texts. The mean and standard deviations of the ASRs are shown in the table. The ‘Overall’ column indicates the average targeted success rate across all tasks. The best performance values for each task are highlighted in **bold**.

Target Prompt	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
unknown	Single-P	0.24 \pm 1.34e-2	0.39 \pm 5.73e-3	0.21 \pm 6.25e-3	0.05 \pm 2.31e-3	0.22 \pm 8.04e-3
	Multi-P	0.67 \pm 7.14e-3	0.86 \pm 2.09e-3	0.64 \pm 1.35e-3	0.31 \pm 1.44e-2	0.62 \pm 8.16e-3
	CroPA	0.92\pm1.07e-2	0.98\pm6.72e-3	0.70\pm3.42e-3	0.34\pm3.19e-3	0.74\pm6.75e-3
I am sorry	Single-P	0.21 \pm 1.50e-3	0.43 \pm 7.52e-3	0.47 \pm 8.59e-3	0.34 \pm 5.01e-3	0.36 \pm 6.28e-3
	Multi-P	0.60 \pm 1.28e-3	0.85 \pm 1.45e-2	0.71 \pm 1.26e-2	0.60 \pm 3.97e-3	0.69 \pm 9.87e-3
	CroPA	0.90\pm3.56e-3	0.96\pm5.25e-3	0.75\pm8.34e-3	0.72\pm7.04e-3	0.83\pm6.31e-3
not sure	Single-P	0.25 \pm 1.42e-3	0.36 \pm 1.52e-3	0.09 \pm 1.25e-2	0.00 \pm 6.04e-3	0.17 \pm 7.03e-3
	Multi-P	0.55 \pm 9.56e-3	0.55 \pm 2.95e-3	0.11 \pm 5.09e-3	0.02 \pm 6.12e-3	0.31 \pm 6.39e-3
	CroPA	0.88\pm1.19e-2	0.86\pm3.79e-3	0.30\pm8.19e-3	0.17\pm9.29e-3	0.55\pm8.82e-3
very good	Single-P	0.35 \pm 8.31e-3	0.52 \pm 1.17e-2	0.15 \pm 4.02e-3	0.05 \pm 9.72e-3	0.27 \pm 8.92e-3
	Multi-P	0.81 \pm 9.51e-3	0.93 \pm 3.38e-3	0.40 \pm 1.91e-3	0.20 \pm 1.42e-2	0.59 \pm 8.79e-2
	CroPA	0.95\pm1.13e-2	0.97\pm5.26e-3	0.64\pm2.36e-3	0.27\pm1.05e-2	0.71\pm8.61e-3
too late	Single-P	0.21 \pm 1.72e-3	0.38 \pm 8.43e-3	0.21 \pm 8.56e-3	0.04 \pm 9.92e-3	0.21 \pm 7.84e-3
	Multi-P	0.78 \pm 2.71e-3	0.90 \pm 7.93e-3	0.54 \pm 1.48e-3	0.17 \pm 1.37e-2	0.60 \pm 8.07e-3
	CroPA	0.90\pm1.03e-2	0.95\pm5.36e-3	0.73\pm8.28e-3	0.20\pm8.65e-3	0.70\pm8.33e-3
metaphor	Single-P	0.26 \pm 1.46e-2	0.56 \pm 8.22e-3	0.50 \pm 5.52e-3	0.14 \pm 1.21e-2	0.37 \pm 8.83e-3
	Multi-P	0.83 \pm 1.46e-2	0.92 \pm 1.18e-2	0.81 \pm 1.41e-2	0.42 \pm 1.35e-2	0.75 \pm 1.36e-2
	CroPA	0.96\pm1.39e-2	0.99\pm2.23e-3	0.92\pm3.74e-3	0.62\pm1.63e-3	0.87\pm1.07e-2

4.3 CROPA MEETS IN-CONTEXT LEARNING

In addition to the textual prompt, the Flamingo model also supports providing extra images as in-context learning examples to improve the task adaptation ability. Whether these in-context learning examples have an influence on the cross-prompt attack remains unclear. Therefore, we tested the ASRs of the image adversarial examples with the number of in-context learning examples different from the one provided in the optimisation stage. During the optimisation stage, the image adversarial examples are updated under the 0-shot setting, namely no extra images are provided as the in-context learning examples. In the evaluations, the 2-shot setting is used, i.e. two extra images are used as the in-context learning examples. Evaluation results under the 0-shot setting are also provided for comparison.

As Table 2 shows, the CroPA still achieves the best performance under the 2-shot settings. We can observe that in-context learning examples can decrease the ASRs, as these two extra in-context learning examples cause a shift in the generation condition different from the optimisation stage.

Table 2: Targeted ASRs of with and without visual in-context learning. The shot indicates the number of images added for in-context learning. The model utilised is Flamingo. The mean and standard deviations of the ASRs are shown in the table. The best performance values for each task are highlighted in **bold**.

Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
Multi-P (shot=0)	0.67 \pm 7.14e-3	0.86 \pm 2.09e-3	0.64 \pm 1.35e-3	0.31 \pm 1.44e-2	0.62 \pm 8.16e-3
CroPA (shot=0)	0.92\pm1.07e-2	0.98\pm6.72e-3	0.70\pm3.42e-3	0.34\pm3.19e-3	0.74\pm6.75e-3
Multi-P (shot=2)	0.59 \pm 6.24e-3	0.81 \pm 1.43e-2	0.50 \pm 1.12e-2	0.25 \pm 9.38e-3	0.54 \pm 3.18e-3
CroPA (shot=2)	0.84\pm3.18e-3	0.96\pm1.18e-3	0.76\pm1.31e-2	0.26\pm9.41e-3	0.70\pm1.09e-2

4.4 CONVERGENCE OF CROPA

In this section, we conduct experiments to compare the performance of baseline and CroPA methods over different update iterations. As shown in Figure 3, we present the results of the overall targeted ASRs with attack iterations from 300 to 1900 every 200 iterations. The number of prompts used for Multi-P and CroPA in optimisation is set to ten.

For all the methods, adding more attack iterations can increase ASRs at the beginning, but the performances eventually converge. For the Single-P method, the improvement in cross-prompt transferability by using more iterations has quickly become marginal after 300 iterations. However, for the Multi-P and CroPA, the performance can still increase after 1000 epochs.




Figure 3: Targeted ASRs of Baseline Methods (Single-P and Multi-P) and CroPA over attack Iterations. With the same number of attack iterations, our CroPA significantly outperforms baselines.

The figure also demonstrated that CroPA methods do not rely on extra attack iterations to gain better performance compared to the baseline approaches. Given the same attack iterations, the CroPA consistently achieves better performance compared to the Single-P and Multi-P methods. CroPA requires only 500 iterations to achieve the target ASR above 0.6 while the stronger baseline Multi-P requires over 1000 iterations.

Overall, Section 4.1 and this section have shown that the cross-prompt transferability of the baseline method is limited compared to CroPA with the increase of the prompt numbers and iterations.

4.5 CROPA WITH DIFFERENT UPDATE STRATEGY

In this section, we explore the effect of different update strategies. As described in previous sections, the update frequency of both the image perturbation and the prompt perturbation can be different. A special case of CroPA is $\text{CroPA}_{\text{joint}}$, where the image perturbation and the prompt perturbation have the same update frequency. We compare the result of CroPA with $\text{CroPA}_{\text{joint}}$ on different tasks with different in-context learning settings. Similar to Section 4.3, adversarial examples are optimised under 0-shot and tested under 0-shot and 2-shot settings.

As shown in Table 3, the CroPA outperform the $\text{CroPA}_{\text{joint}}$ overall and most individual tasks with different in-context learning image examples. The stronger performance of CroPA derives from its flexibility in choosing the update step size of the prompt perturbation. As presented in Appendix A, the $\text{CroPA}_{\text{joint}}$ is sensitive to the step size of the prompt embedding: if the prompt update size is too large the optimisation fails to converge. CroPA is more tolerant of large prompt update sizes by reducing the prompt update frequency.

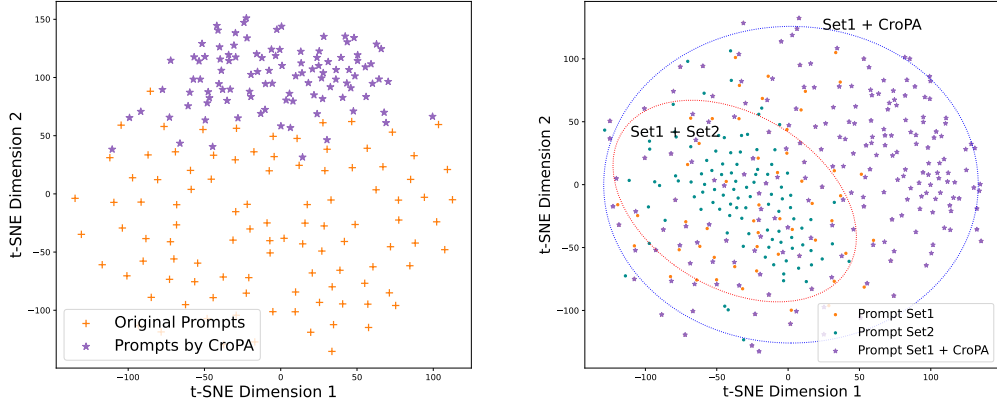
Table 3: Our CroPA with alternative optimization outperforms the one with vision and prompt joint optimization in most cases. The mean and standard deviations of the ASRs are shown in the table. The best performance values for each task are highlighted in **bold**

Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
$\text{CroPA}_{\text{joint}}$ (shot=0)	$0.86 \pm 1.44e-3$	$0.95 \pm 9.91e-3$	$0.73 \pm 5.40e-3$	$0.31 \pm 8.11e-3$	$0.71 \pm 6.99e-3$
CroPA (shot=0)	$0.92 \pm 1.07e-2$	$0.98 \pm 6.72e-3$	$0.70 \pm 3.42e-3$	$0.34 \pm 3.19e-3$	$0.74 \pm 6.75e-3$
$\text{CroPA}_{\text{joint}}$ (shot=2)	$0.76 \pm 4.49e-3$	$0.94 \pm 6.74e-3$	$0.71 \pm 1.15e-2$	$0.25 \pm 4.20e-3$	$0.66 \pm 7.37e-3$
CroPA (shot=2)	$0.84 \pm 3.18e-3$	$0.96 \pm 1.18e-3$	$0.76 \pm 1.31e-2$	$0.26 \pm 9.41e-3$	$0.70 \pm 1.09e-2$

4.6 UNDERSTANDING THE EFFECTIVENESS OF CROPA METHODS

Visualisation of the Prompt Embedding Coverage To explore the underlying reasons for the better performance of CroPA compared to the baseline approach, we visualise the sentence embedding of the original prompt and perturbed prompts by CroPA, which is obtained by the averaging embedding of each token.

As demonstrated in Figure 4a, the orange plus symbol denotes the original prompts while the purple star symbol denotes the original embedding added with the perturbation δ_t . It can be observed that there is almost no overlap between the prompt embedding perturbed by CroPA and the original prompt embedding. This verifies that the adversarial prompt effectively increases the coverage of the original embedding.



(a) Visualisation of the prompt embeddings and its prompt embedding created by CroPA. The orange plus symbol denotes the original prompt and the purple star symbol denotes the embedding by adding the adversarial prompt perturbation to the original prompt embedding.

(b) Comparison between the embedding coverage difference between prompts generated by CroPA and extra Prompt Set 2. The red circle denote the embedding coverage of Prompt Set 1 and Prompt Set 2, while the blue circle represents the coverage of Prompt Set 1 with the prompts generated by CroPA.

Figure 4: Visualisation of the prompt embeddings with t-SNE (Van der Maaten & Hinton, 2008).

To have a clearer comparison between the baseline approach, which relies on simply adding more prompts, and the CroPA methods, we visualise the coverage of prompt embeddings of these two methods in Figure 4b shows. The embedding of Prompt Set 1 is denoted by the orange dots, while the embeddings of Prompt Set 2 are denoted as the cyan dots. In the CroPA, only Prompt Set 1 is provided. By introducing the learnable prompt perturbation to Prompts Set 1, the prompt embeddings that have been covered during optimisation are denoted by the purple stars. The blue eclipse is the approximated coverage of the prompt embedding for CroPA using Prompt Set 1. For the baseline method, both Prompt Set 1 and Prompt Set 2 are provided and the red eclipse approximately represents the coverage of their coverage. It can be observed that with only Prompt Set 1, the area covered by the CroPA is broader than the one covered by the embeddings of Prompt Set 1 and Prompt Set 2.

The visualisation of difference in the prompt embedding coverage explains the reason why the CroPA methods can outperform the baseline approach even if the number of prompts used in optimisation is less than the baseline approach.

Prompt Embedding Decoding We explored the decoding of the adversarial prompt embedding to a human-readable text format. The embedding of each token is decoded to the readable text closest in terms of cosine distance using the pre-trained embedding look-up table of the language models. The results show that all the perturbed embeddings are still closest to their original tokens. This finding also supports the effectiveness of CroPA: There exists prompt embedding that cannot be represented by human-readable text. Therefore, even if given a sufficient number of prompts in the baseline approach, it still can not cover all the prompt embedding space of the adversarial prompt in the CroPA framework.

5 CONCLUSION

In this paper, we first raise an interesting and important question, can a single adversarial example mislead all predictions of a vision-language model with different prompts? We formulate the essence of the question as the cross-prompt adversarial transferability of adversarial perturbation. Extensive experiments show that intuitive baseline approaches only achieve limited transferability and our proposed CroPA improves the transferability significantly on different VLMs in different multi-modal tasks. One of the ways to further improve the practical applicability of our method is to implement the optimization with query-based strategies (Chen et al., 2017b; Ilyas et al., 2018), which we leave to future work.

Acknowledgement This work is supported by the UKRI grant: Turing AI Fellowship EP/W002981/1, and EPSRC/MURI grant: EP/N019474/1. We would also like to thank the Royal Academy of Engineering and FiveAI.

REFERENCES

- Nayyer Aafaq, Naveed Akhtar, Wei Liu, Mubarak Shah, and Ajmal Mian. Controlled caption generation for images through adversarial attacks. *arXiv preprint arXiv:2107.03050*, 2021.
- Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katherine Millican, Malcolm Reynolds, et al. Flamingo: a visual language model for few-shot learning. *Advances in Neural Information Processing Systems*, 35:23716–23736, 2022.
- Anas Awadalla, Irena Gao, Josh Gardner, Jack Hessel, Yusuf Hanafy, Wanrong Zhu, Kalyani Marathe, Yonatan Bitton, Samir Gadre, Shiori Sagawa, et al. Openflamingo: An open-source framework for training large autoregressive vision-language models. *arXiv preprint arXiv:2308.01390*, 2023.
- Hongge Chen, Huan Zhang, Pin-Yu Chen, Jinfeng Yi, and Cho-Jui Hsieh. Attacking visual language grounding with adversarial examples: A case study on neural image captioning. *arXiv preprint arXiv:1712.02051*, 2017a.
- Pin-Yu Chen, Huan Zhang, Yash Sharma, Jinfeng Yi, and Cho-Jui Hsieh. Zoo: Zeroth order optimization based black-box attacks to deep neural networks without training substitute models. In *Proceedings of the 10th ACM workshop on artificial intelligence and security*, pp. 15–26, 2017b.
- Wenliang Dai, Junnan Li, Dongxu Li, Anthony Meng Huat Tiong, Junqi Zhao, Weisheng Wang, Boyang Li, Pascale Fung, and Steven C. H. Hoi. Instructblip: Towards general-purpose vision-language models with instruction tuning. *ArXiv*, abs/2305.06500, 2023. URL <https://api.semanticscholar.org/CorpusID:258615266>.
- Ian J Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*, 2014.
- Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 6904–6913, 2017.
- Jindong Gu, Hengshuang Zhao, Volker Tresp, and Philip HS Torr. Segpgd: An effective and efficient adversarial attack for evaluating and boosting segmentation robustness. In *European Conference on Computer Vision*, pp. 308–325. Springer, 2022.
- Jindong Gu, Zhen Han, Shuo Chen, Ahmad Beirami, Bailan He, Gengyuan Zhang, Ruotong Liao, Yao Qin, Volker Tresp, and Philip Torr. A systematic survey of prompt engineering on vision-language foundation models. *arXiv preprint arXiv:2307.12980*, 2023a.
- Jindong Gu, Xiaojun Jia, Pau de Jorge, Wenqian Yu, Xinwei Liu, Avery Ma, Yuan Xun, Anjun Hu, Ashkan Khakzar, Zhijiang Li, et al. A survey on transferability of adversarial examples across deep neural networks. *arXiv preprint arXiv:2310.17626*, 2023b.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Andrew Ilyas, Logan Engstrom, Anish Athalye, and Jessy Lin. Black-box adversarial attacks with limited queries and information. In *International conference on machine learning*, pp. 2137–2146. PMLR, 2018.
- Divyansh Kaushik, Douwe Kiela, Zachary C Lipton, and Wen-tau Yih. On the efficacy of adversarial data collection for question answering: Results from a large-scale randomized study. *arXiv preprint arXiv:2106.00872*, 2021.

- Venelin Kovatchev, Trina Chatterjee, Venkata S Govindarajan, Jifan Chen, Eunsol Choi, Gabriella Chronis, Anubrata Das, Katrin Erk, Matthew Lease, Junyi Jessy Li, et al. longhorns at dadc 2022: How many linguists does it take to fool a question answering model? a systematic approach to adversarial attacks. *arXiv preprint arXiv:2206.14729*, 2022.
- Junnan Li, Dongxu Li, Caiming Xiong, and Steven Hoi. Blip: Bootstrapping language-image pre-training for unified vision-language understanding and generation. In *International Conference on Machine Learning*, pp. 12888–12900. PMLR, 2022.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. *arXiv preprint arXiv:2301.12597*, 2023.
- Linjie Li, Jie Lei, Zhe Gan, and Jingjing Liu. Adversarial vqa: A new benchmark for evaluating the robustness of vqa models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 2042–2051, 2021.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer Vision—ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6–12, 2014, Proceedings, Part V 13*, pp. 740–755. Springer, 2014.
- Yanpei Liu, Xinyun Chen, Chang Liu, and Dawn Song. Delving into transferable adversarial examples and black-box attacks. *arXiv preprint arXiv:1611.02770*, 2016.
- Yantao Lu, Yunhan Jia, Jianyu Wang, Bai Li, Weiheng Chai, Lawrence Carin, and Senem Velipasalar. Enhancing cross-task black-box transferability of adversarial examples with dispersion reduction. In *Proceedings of the IEEE/CVF conference on Computer Vision and Pattern Recognition*, pp. 940–949, 2020.
- Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. Towards deep learning models resistant to adversarial attacks. *arXiv preprint arXiv:1706.06083*, 2017.
- Seyed-Mohsen Moosavi-Dezfooli, Alhussein Fawzi, Omar Fawzi, and Pascal Frossard. Universal adversarial perturbations. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 1765–1773, 2017.
- Konda Reddy Mopuri, Utsav Garg, and R Venkatesh Babu. Fast feature fool: A data independent approach to universal adversarial perturbations. *arXiv preprint arXiv:1707.05572*, 2017.
- Muhammad Muzammal Naseer, Salman H Khan, Muhammad Haris Khan, Fahad Shahbaz Khan, and Fatih Porikli. Cross-domain transferability of adversarial perturbations. *Advances in Neural Information Processing Systems*, 32, 2019.
- Muzammal Naseer, Salman H Khan, Shafin Rahman, and Fatih Porikli. Task-generalizable adversarial attack based on perceptual metric. *arXiv preprint arXiv:1811.09020*, 2018.
- Nicolas Papernot, Patrick McDaniel, and Ian Goodfellow. Transferability in machine learning: from phenomena to black-box attacks using adversarial samples. *arXiv preprint arXiv:1605.07277*, 2016.
- Mathieu Salzmann et al. Learning transferable adversarial perturbations. *Advances in Neural Information Processing Systems*, 34:13950–13962, 2021.
- Sasha Sheng, Amanpreet Singh, Vedanuj Goswami, Jose Magana, Tristan Thrush, Wojciech Galuba, Devi Parikh, and Douwe Kiela. Human-adversarial visual question answering. *Advances in Neural Information Processing Systems*, 34:20346–20359, 2021.
- Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. Intriguing properties of neural networks. In *arXiv preprint arXiv:1312.6199*, 2013.


- Florian Tramèr, Nicolas Papernot, Ian Goodfellow, Dan Boneh, and Patrick McDaniel. The space of transferable adversarial examples. *arXiv preprint arXiv:1704.03453*, 2017.
- Laurens Van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine learning research*, 9(11), 2008.
- Boxi Wu, Jindong Gu, Zhifeng Li, Deng Cai, Xiaofei He, and Wei Liu. Towards efficient adversarial training on vision transformers. In *European Conference on Computer Vision*, pp. 307–325. Springer, 2022.
- Xiaojun Xu, Xinyun Chen, Chang Liu, Anna Rohrbach, Trevor Darrell, and Dawn Song. Fooling vision and language models despite localization and attention mechanism. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 4951–4961, 2018.
- Yan Xu, Baoyuan Wu, Fumin Shen, Yanbo Fan, Yong Zhang, Heng Tao Shen, and Wei Liu. Exact adversarial attack to image captioning via structured output learning with latent variables. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4135–4144, 2019.
- Xu Yang, Kaihua Tang, Hanwang Zhang, and Jianfei Cai. Auto-encoding scene graphs for image captioning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 10685–10694, 2019.
- Ting Yao, Yingwei Pan, Yehao Li, and Tao Mei. Exploring visual relationship for image captioning. In *Proceedings of the European conference on computer vision (ECCV)*, pp. 684–699, 2018.
- Wenqian Yu, Jindong Gu, Zhijiang Li, and Philip Torr. Reliable evaluation of adversarial transferability. *arXiv preprint arXiv:2306.08565*, 2023.
- Jiaming Zhang, Qi Yi, and Jitao Sang. Towards adversarial attack on vision-language pre-training models. In *Proceedings of the 30th ACM International Conference on Multimedia*, pp. 5005–5013, 2022a.
- Shaofeng Zhang, Zheng Wang, Xing Xu, Xiang Guan, and Yang Yang. Fooled by imagination: Adversarial attack to image captioning via perturbation in complex domain. In *2020 IEEE International Conference on Multimedia and Expo (ICME)*, pp. 1–6. IEEE, 2020.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022b.
- Yunqing Zhao, Tianyu Pang, Chao Du, Xiao Yang, Chongxuan Li, Ngai-Man Cheung, and Min Lin. On evaluating adversarial robustness of large vision-language models. *arXiv preprint arXiv:2305.16934*, 2023.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *arXiv preprint arXiv:2306.05685*, 2023.
- Deyao Zhu, Jun Chen, Xiaoqian Shen, Xiang Li, and Mohamed Elhoseiny. Minigpt-4: Enhancing vision-language understanding with advanced large language models. *arXiv preprint arXiv:2304.10592*, 2023.
- Andy Zou, Zifan Wang, J Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

A EFFECT OF PROMPT PERTURBATION UPDATE STEP SIZE

Selecting the correct prompt updating step size is vital for the CroPA_{joint} variant. We present the results utilising the validation dataset to probe the optimal prompt update step size. The image adversarial update step size is fixed to 1/255. The number of prompts during optimisation is set to ten. The update step size used in this experiment are 0.01, 0.001, 0.0001, 1e-05 and 1e-06.

As depicted in Figure 5a, we can observe that if the update step size is set to 0.01, though the loss decreases during the first one hundred iterations, it fails to converge eventually. The corresponding performance recorded in Figure 5b validates this, it achieves the lowest targeted ASR compared to the other step size.

For the rest of the learning rate, we can observe that the optimisation processes are roughly converged though the values it converge to can be different. It is notable that the update step size 0.001, which is larger than the rest of the update step size, achieves the best performance even if the final loss value it converges is not the smallest compared to the other prompt update step size.



(a) The optimisation loss of the CroPA_{joint} algorithm with different update step sizes of the prompts. The image adversarial update step size is fixed to 1/255. (b) The targeted attack success rate of the CroPA_{joint} algorithm with different update step sizes of the prompts.

Figure 5: Effect of different update step size of prompt perturbation of CroPA_{joint}

Also, we can notice that at around five hundred iterations, the optimisation loss increases slightly, but the performance of the targeted success rate increases significantly. This is because a larger step size of the prompt can be perceived as a stronger adversary for the image adversarial perturbation. During this learning process, the image adversarial perturbation is optimised to counteract this challenging prompt input and thereby obtains stronger cross-prompt adversarial transferability. This empirical finding is consistent with the idea in machine learning that the models are optimised on challenging cases to improve the model’s performance. This finding is also verified from the other perspective. We can observe that the performance with the step size set to 1e-06 increases most slowly compared to the performance with a larger update step size of prompts.

In conclusion, the optimal choice of the prompt update step size should allow the image perturbation to converge at a reasonably large value range to allow the perturbation gain stronger cross-prompt adversarial transferability.

B PROMPTS FOR DIFFERENT TASKS

Prompts for VQA

Any cutlery items visible in the image?
Any bicycles visible in this image?
Any boats visible in the image?
Any bottles present in the image?
Are curtains noticeable in the image?
Are flags present in the image?
Are flowers present in the image?
Are fruits present in the image?

Are glasses discernible in the image?
Are hills visible in the image?
Are plates discernible in the image?
Are shoes visible in this image?
Are there any insects in the image?
Are there any ladders in the image?
Are there any man-made structures in the image?
Are there any signs or markings in the image?
Are there any street signs in the image?
Are there balloons in the image?
Are there bridges in the image?
Are there musical notes in the image?
Are there people sitting in the image?
Are there skyscrapers in the image?
Are there toys in the image?
Are toys present in this image?
Are umbrellas discernible in the image?
Are windows visible in the image?
Can birds be seen in this image?
Can stars be seen in this image?
Can we find any bags in this image?
Can you find a crowd in the image?
Can you find a hat in the image?
Can you find any musical instruments in this image?
Can you identify a clock in this image?
Can you identify a computer in this image?
Can you see a beach in the image?
Can you see a bus in the image?
Can you see a mailbox in the image?
Can you see a mountain in the image?
Can you see a staircase in the image?
Can you see a stove or oven in the image?
Can you see a sunset in the image?
Can you see any cups or mugs in the image?
Can you see any jewelry in the image?
Can you see shadows in the image?
Can you see the sky in the image?
Can you spot a candle in this image?
Can you spot a farm in this image?
Can you spot a pair of shoes in the image?
Can you spot a rug or carpet in the image?
Can you spot any dogs in the image?
Can you spot any snow in the image?
Do you notice a bicycle in the image?
Does a ball feature in this image?
Does a bridge appear in the image?
Does a cat appear in the image?
Does a fence appear in the image?
Does a fire feature in this image?
Does a mirror feature in this image?
Does a table feature in this image?
Does it appear to be nighttime in the image?
Does it look like an outdoor image?
Does it seem to be countryside in the image?
Does the image appear to be a cartoon or comic strip?
Does the image contain any books?
Does the image contain any electronic devices?
Does the image depict a road?
Does the image display a river?

Does the image display any towers?
Does the image feature any art pieces?
Does the image have a lamp?
Does the image have any pillows?
Does the image have any vehicles?
Does the image have furniture?
Does the image primarily display natural elements?
Does the image seem like it was taken during the day?
Does the image seem to be taken indoors?
Does the image show any airplanes?
Does the image show any benches?
Does the image show any landscapes?
Does the image show any movement?
Does the image show any sculptures?
Does the image show any signs?
Does the image show food?
Does the image showcase a building?
How many animals are present in the image?
How many bikes are present in the image?
How many birds are visible in the image?
How many buildings can be identified in the image?
How many cars can be seen in the image?
How many doors can you spot in the image?
How many flowers can be identified in the image?
How many trees feature in the image?
Is a chair noticeable in the image?
Is a computer visible in the image?
Is a forest noticeable in the image?
Is a painting visible in the image?
Is a path or trail visible in the image?
Is a phone discernible in the image?
Is a train noticeable in the image?
Is sand visible in the image?
Is the image displaying any clouds?
Is the image set in a city environment?
Is there a plant in the image?
Is there a source of light visible in the image?
Is there a television displayed in the image?
Is there grass in the image?
Is there text in the image?
Is water visible in the image, like a sea, lake, or river?
How many people are captured in the image?
How many windows can you count in the image?
How many animals, other than birds, are present?
How many statues or monuments stand prominently in the scene?
How many streetlights are visible?
How many items of clothing can you identify?
How many shoes can be seen in the image?
How many clouds appear in the sky?
How many pathways or trails are evident?
How many bridges can you spot?
How many boats are present, if it's a waterscape?
How many pieces of fruit can you identify?
How many hats are being worn by people?
How many different textures can you discern?
How many signs or billboards are visible?
How many musical instruments can be seen?
How many flags are present in the image?
How many mountains or hills can you identify?

How many books are visible, if any?
How many bodies of water, like ponds or pools, are in the scene?
How many shadows can you spot?
How many handheld devices, like phones, are present?
How many pieces of jewelry can be identified?
How many reflections, perhaps in mirrors or water, are evident?
How many pieces of artwork or sculptures can you see?
How many staircases or steps are in the image?
How many archways or tunnels can be counted?
How many tools or equipment are visible?
How many modes of transportation, other than cars and bikes, can you spot?
How many lamp posts or light sources are there?
How many plants, other than trees and flowers, feature in the scene?
How many fences or barriers can be seen?
How many chairs or seating arrangements can you identify?
How many different patterns or motifs are evident in clothing or objects?
How many dishes or food items are visible on a table setting?
How many glasses or mugs can you spot?
How many pets or domestic animals are in the scene?
How many electronic gadgets can be counted?
Where is the brightest point in the image?
Where are the darkest areas located?
Where can one find leading lines directing the viewer's eyes?
Where is the visual center of gravity in the image?
Where are the primary and secondary subjects positioned?
Where do the most vibrant colors appear?
Where is the most contrasting part of the image located?
Where does the image place emphasis through scale or size?
Where do the textures in the image change or transition?
Where does the image break traditional compositional rules?
Where do you see repetition or patterns emerging?
Where does the image exhibit depth or layers?
Where are the boundary lines or borders in the image?
Where do different elements in the image intersect or overlap?
Where does the image hint at motion or movement?
Where are the calm or restful areas of the image?
Where does the image become abstract or less defined?
Where do you see reflections, be it in water, glass, or other surfaces?
Where does the image provide contextual clues about its setting?
Where are the most detailed parts of the image?
Where do you see shadows, and how do they impact the composition?
Where can you identify different geometric shapes?
Where does the image appear to have been cropped or framed intentionally?
Where do you see harmony or unity among the elements?
Where are there disruptions or interruptions in patterns?
What is the spacing between objects or subjects in the image?
What foreground, mid-ground, and background elements can be differentiated?
What type of energy or vibe does the image exude?
What might be the sound environment based on the image's content?
What abstract ideas or concepts does the image seem to touch upon?
What is the relationship between the main subjects in the image?
What items in the image could be considered rare or unique?
What is the gradient or transition of colors like in the image?
What might be the smell or aroma based on the image's content?
What type of textures can be felt if one could touch the image's content?
What boundaries or limits are depicted in the image?
What is the socioeconomic context implied by the image?
What might be the immediate aftermath of the scene in the image?
What seems to be the main source of tension or harmony in the image?

What might be the narrative or backstory of the main subject?
What elements of the image give it its primary visual weight?
Would you describe the image as bright or dark?
Would you describe the image as colorful or dull?

Prompts for Image Classification

Identify the primary theme of this image in one word.
How would you label this image with a single descriptor?
Determine the main category for this image.
Offer a one-word identifier for this picture.
If this image were a file on your computer, what would its name be?
Tag this image with its most relevant keyword.
Provide the primary classification for this photograph.
How would you succinctly categorize this image?
Offer the primary descriptor for the content of this image.
If this image were a product, what label would you place on its box?
Choose a single word that encapsulates the image's content.
How would you classify this image in a database?
In one word, describe the essence of this image.
Provide the most fitting category for this image.
What is the principal subject of this image?
If this image were in a store, which aisle would it belong to?
Provide a singular term that characterizes this picture.
How would you caption this image in a photo contest?
Select a label that fits the main theme of this image.
Offer the most appropriate tag for this image.
Which keyword best summarizes this image?
How would you title this image in an exhibition?
Provide a succinct identifier for the image's content.
Choose a word that best groups this image with others like it.
If this image were in a museum, how would it be labeled?
Assign a central theme to this image in one word.
Tag this photograph with its primary descriptor.
What is the overriding theme of this picture?
Provide a classification term for this image.
How would you sort this image in a collection?
Identify the main subject of this image concisely.
If this image were a magazine cover, what would its title be?
What term would you use to catalog this image?
Classify this picture with a singular term.
If this image were a chapter in a book, what would its title be?
Select the most fitting classification for this image.
Define the essence of this image in one word.
How would you label this image for easy retrieval?
Determine the core theme of this photograph.
In a word, encapsulate the main subject of this image.
If this image were an art piece, how would it be labeled in a gallery?
Provide the most concise descriptor for this picture.
How would you name this image in a photo archive?
Choose a word that defines the image's main content.
What would be the header for this image in a catalog?
Classify the primary essence of this picture.
What label would best fit this image in a slideshow?
Determine the dominant category for this photograph.
Offer the core descriptor for this image.
If this image were in a textbook, how would it be labeled in the index?
Select the keyword that best defines this image's theme.

Provide a classification label for this image.
If this image were a song title, what would it be?
Identify the main genre of this picture.
Assign the most apt category to this image.
Describe the overarching theme of this image in one word.
What descriptor would you use for this image in a portfolio?
Summarize the image's content with a single identifier.
Imagine you're explaining this image to someone over the phone. Please describe the image in one word?
Perform the image classification task on this image. Give the label in one word.
Imagine a child is trying to identify the image. What might they excitedly point to and name?
If this image were turned into a jigsaw puzzle, what would the box label say to describe the picture inside?
Classify the content of this image.
If you were to label this image, what label would you give?
What category best describes this image?
Describe the central subject of this image in a single word.
Provide a classification for the object depicted in this image.
If this image were in a photo album, what would its label be?
Categorize the content of the image.
If you were to sort this image into a category, which one would it be?
What keyword would you associate with this image?
Assign a relevant classification to this image.
If this image were in a gallery, under which section would it belong?
Describe the main theme of this image in one word.
Under which category would this image be cataloged in a library?
What classification tag fits this image the best?
Provide a one-word description of this image's content.
If you were to archive this image, what descriptor would you use?

Prompts for Image Captioning

Elaborate on the elements present in this image.
In one sentence, summarize the activity in this image.
Relate the main components of this picture in words.
What narrative unfolds in this image?
Break down the main subjects of this photo.
Give an account of the main scene in this image.
In a few words, state what this image represents.
Describe the setting or location captured in this photograph.
Provide an overview of the subjects or objects seen in this picture.
Identify the primary focus or point of interest in this image.
What would be the perfect title for this image?
How would you introduce this image in a presentation?
Present a quick rundown of the image's main subject.
What's the key event or subject captured in this photograph?
Relate the actions or events taking place in this image.
Convey the content of this photograph in a single phrase.
Offer a succinct description of this picture.
Give a concise overview of this image.
Translate the contents of this picture into a sentence.
Describe the characters or subjects seen in this image.
Capture the activities happening in this image with words.
How would you introduce this image to an audience?
State the primary events or subjects in this picture.
What are the main elements in this photograph?
Provide an interpretation of this image's main event or subject.
How would you title this image for an art gallery?

What scenario or setting is depicted in this image?
Concisely state the main actions occurring in this image.
Offer a short summary of this photograph's contents.
How would you annotate this image in an album?
If you were to describe this image on the radio, how would you do it?
In your own words, narrate the main event in this image.
What are the notable features of this image?
Break down the story this image is trying to tell.
Describe the environment or backdrop in this photograph.
How would you label this image in a catalog?
Convey the main theme of this picture succinctly.
Characterize the primary event or action in this image.
Provide a concise depiction of this photo's content.
Write a brief overview of what's taking place in this image.
Illustrate the main theme of this image with words.
How would you describe this image in a gallery exhibit?
Highlight the central subjects or actions in this image.
Offer a brief narrative of the events in this photograph.
Translate the activities in this image into a brief sentence.
Give a quick rundown of the primary subjects in this image.
Provide a quick summary of the scene captured in this photo.
How would you explain this image to a child?
What are the dominant subjects or objects in this photograph?
Summarize the main events or actions in this image.
Describe the context or setting of this image briefly.
Offer a short description of the subjects present in this image.
Detail the main scenario or setting seen in this picture.
Describe the main activities or events unfolding in this image.
Provide a concise explanation of the content in this image.
If this image were in a textbook, how would it be captioned?
Provide a summary of the primary focus of this image.
State the narrative or story portrayed in this picture.
How would you introduce this image in a documentary?
Detail the subjects or events captured in this image.
Offer a brief account of the scenario depicted in this photograph.
State the main elements present in this image concisely.
Describe the actions or events happening in this picture.
Provide a snapshot description of this image's content.
How would you briefly describe this image's main subject or event?
Describe the content of this image.
What's happening in this image?
Provide a brief caption for this image.
Tell a story about this image in one sentence.
If this image could speak, what would it say?
Summarize the scenario depicted in this image.
What is the central theme or event shown in the picture?
Create a headline for this image.
Explain the scene captured in this image.
If this were a postcard, what message would it convey?
Narrate the visual elements present in this image.
Give a short title to this image.
How would you describe this image to someone who can't see it?
Detail the primary action or subject in the photo.
If this image were the cover of a book, what would its title be?
Translate the emotion or event of this image into words.
Compose a one-liner describing this image's content.
Imagine this image in a magazine. What caption would go with it?
Capture the essence of this image in a brief description.

Narrate the visual story displayed in this photograph.

C DETAILED DATA

Table 4: Targeted attack success rates tested on **Flamingo**. The columns labeled ‘VQA_{general}’ and ‘VQA_{specific}’ denote the VQA prompts for general and specific types of questions respectively. The columns ‘Classification’ and ‘Captioning’ refer to the success rates for the image classification and captioning tasks respectively. The ‘Overall’ column indicates the average targeted success rate across all tasks. The ‘Num’ column signifies the number of prompts used during optimisation. ‘CroPA_{joint}’ and ‘CroPA’ represent the CroPA methods utilizing joint and alternating updating strategies respectively. The best performance values for each task are highlighted in bold.

No. of Prompts	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
1	Single-P	0.24	0.39	0.21	0.05	0.22
	CroPA _{joint}	0.30	0.47	0.16	0.09	0.26
	CroPA	0.52	0.69	0.38	0.17	0.44
5	Baseline	0.63	0.82	0.57	0.22	0.56
	CroPA _{joint}	0.82	0.93	0.69	0.30	0.69
	CroPA	0.90	0.96	0.56	0.39	0.70
10	Baseline	0.67	0.86	0.64	0.31	0.62
	CroPA _{joint}	0.86	0.95	0.73	0.31	0.71
	CroPA	0.92	0.98	0.70	0.34	0.74
50	Baseline	0.67	0.85	0.50	0.25	0.57
	CroPA _{joint}	0.88	0.94	0.73	0.33	0.72
	CroPA	0.95	0.99	0.67	0.40	0.75
100	Baseline	0.70	0.85	0.57	0.29	0.60
	CroPA _{joint}	0.90	0.95	0.74	0.35	0.74
	CroPA	0.96	0.99	0.68	0.44	0.77

Table 5: Targeted attack success rates tested on **BLIP-2**. The best performance values for each task are highlighted in bold.

No. of Prompts	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
1	Baseline	0.24	0.34	0.45	0.32	0.34
	CroPA _{joint}	0.41	0.48	0.48	0.41	0.45
	CroPA	0.52	0.63	0.65	0.58	0.60
5	Baseline	0.51	0.59	0.62	0.58	0.58
	CroPA _{joint}	0.80	0.83	0.75	0.81	0.80
	CroPA	0.81	0.83	0.80	0.84	0.82
10	Baseline	0.68	0.81	0.68	0.67	0.71
	CroPA _{joint}	0.83	0.83	0.75	0.79	0.80
	CroPA	0.86	0.90	0.82	0.84	0.86
50	Baseline	0.67	0.74	0.67	0.72	0.70
	CroPA _{joint}	0.84	0.88	0.79	0.84	0.84
	CroPA	0.90	0.93	0.87	0.91	0.90
100	Baseline	0.67	0.76	0.68	0.66	0.69
	CroPA _{joint}	0.87	0.93	0.81	0.87	0.87
	CroPA	0.95	0.95	0.87	0.92	0.92

D CROPA WITH NON-TARGETED ATTACK GOAL

As presented in Table 11, we tested the effectiveness of our approach under the non-target attack setting. In this setting, the attack is considered to be successful if the model is misled to produce any different predictions. The CroPA method also demonstrated superior performance compared to the baseline approach under different numbers of prompts used in the optimisation.

Table 6: Targeted attack success rates with RandomRotation as the defense strategy. During the optimisation phase, no additional data augmentation is used. The best performance values for each task are highlighted.

Setting	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
Without Defense	Multi-P	0.67	0.86	0.64	0.31	0.62
	CroPA	0.92	0.98	0.70	0.34	0.74
With Defense	Multi-P	0.58	0.79	0.52	0.26	0.54
	CroPA	0.89	0.95	0.61	0.34	0.70

Table 7: The targeted ASR results tested on Flamingo given different perturbation sizes: 8/255, 16/255 and 32/255. The best performance values for each task are highlighted in bold.

Perturbation size	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
8/255	Multi-P	0.37	0.59	0.56	0.06	0.39
	CroPA	0.53	0.75	0.48	0.03	0.45
16/255	Multi-P	0.67	0.86	0.64	0.31	0.62
	CroPA	0.92	0.98	0.70	0.34	0.74
32/255	Multi-P	0.85	0.95	0.47	0.32	0.64
	CroPA	0.98	0.99	0.59	0.42	0.75

Table 8: The cross model test under different settings. "BLIP2 to InstructBLIP" means the perturbations optimised on BLIP2 are tested on the InstructBLIP and "InstructBLIP to BLIP2" are different. The language model for BLIP2 is OPT-2.7b and the model for InstructBLIP is Vicuna-7b. The best performance values for each task are highlighted in bold.

Settings	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
BLIP2 to InstructBLIP	Multi-P	0.00	0.01	0.04	0.03	0.02
	CroPA	0.00	0.04	0.15	0.11	0.08
InstructBLIP to BLIP2	Multi-P	0.00	0.02	0.10	0.02	0.04
	CroPA	0.01	0.05	0.13	0.04	0.06

Table 9: The targeted ASR results when the target texts are set to harmful instructions. The best performance values for each task are highlighted in bold.

Target Texts	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
suicide	Multi-P	0.69	0.83	0.75	0.61	0.72
	CroPA	0.84	0.91	0.91	0.78	0.86
kidnap	Multi-P	0.87	0.92	0.69	0.73	0.80
	CroPA	0.94	0.96	0.88	0.81	0.90
bomb	Multi-P	0.68	0.82	0.90	0.60	0.75
	CroPA	0.80	0.90	0.94	0.70	0.84

Table 10: The targeted ASR results when the target texts are longer. The best performance values for each task are highlighted in bold.

Target Texts	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
I am sorry	Multi-P	0.60	0.85	0.71	0.60	0.69
	CroPA	0.90	0.96	0.75	0.72	0.83
I cannot answer	Multi-P	0.40	0.66	0.33	0.07	0.37
	CroPA	0.58	0.67	0.33	0.20	0.45
I do not know	Multi-P	0.67	0.75	0.41	0.03	0.47
	CroPA	0.70	0.80	0.43	0.04	0.49
I need a new phone	Multi-P	0.68	0.86	0.85	0.53	0.73
	CroPA	0.83	0.85	0.77	0.70	0.79

Table 11: Non-targeted ASRs of the baseline and CroPA method when different number of prompts are presented. The mean and standard deviations of the ASRs are shown in the table. The best performance values for each task are highlighted in **bold**.

No. of Prompts	Method	VQA _{general}	VQA _{specific}	Classification	Captioning	Overall
1	Single-P	$0.48 \pm 1.21e-2$	$0.63 \pm 8.03e-3$	$0.58 \pm 9.07e-3$	$0.76 \pm 7.89e-3$	$0.61 \pm 9.43e-3$
	CroPA	$0.50 \pm 1.11e-2$	$0.64 \pm 4.93e-3$	$0.64 \pm 1.34e-3$	$0.74 \pm 1.00e-2$	$0.63 \pm 7.91e-3$
5	Multi-P	$0.50 \pm 1.41e-2$	$0.69 \pm 1.43e-2$	$0.70 \pm 1.38e-2$	$0.74 \pm 6.18e-3$	$0.66 \pm 1.26e-2$
	CroPA	$0.64 \pm 1.39e-2$	$0.75 \pm 6.99e-3$	$0.76 \pm 1.29e-2$	$0.79 \pm 5.43e-3$	$0.74 \pm 1.29e-2$
10	Multi-P	$0.49 \pm 5.12e-3$	$0.70 \pm 6.39e-3$	$0.69 \pm 1.13e-2$	$0.76 \pm 5.43e-3$	$0.66 \pm 8.11e-3$
	CroPA	$0.62 \pm 8.79e-3$	$0.77 \pm 1.41e-2$	$0.77 \pm 1.07e-2$	$0.79 \pm 8.98e-3$	$0.74 \pm 1.08e-2$
50	Multi-P	$0.52 \pm 9.61e-3$	$0.72 \pm 1.48e-2$	$0.68 \pm 2.96e-3$	$0.73 \pm 8.25e-3$	$0.66 \pm 9.87e-3$
	CroPA	$0.68 \pm 1.13e-2$	$0.84 \pm 1.07e-2$	$0.81 \pm 1.08e-2$	$0.84 \pm 6.03e-3$	$0.79 \pm 9.98e-3$
100	Multi-P	$0.51 \pm 1.23e-2$	$0.72 \pm 1.23e-2$	$0.73 \pm 1.31e-2$	$0.76 \pm 1.34e-2$	$0.68 \pm 1.29e-2$
	CroPA	$0.69 \pm 8.02e-3$	$0.84 \pm 1.22e-2$	$0.89 \pm 1.01e-2$	$0.80 \pm 1.08e-2$	$0.81 \pm 1.03e-2$